

Project 2

New Attempt

Due Oct 7 by 11:59pm **Points** 100 **Submitting** a file upload

Matrix Multiplication on Map-Reduce

Description

The purpose of this project is to develop a Map-Reduce program on Hadoop to multiply two sparse matrices.

This project must be done individually. No copying is permitted. **Note: We will use a system for detecting software plagiarism, called Moss (<http://theory.stanford.edu/~aiken/moss/>), which is an automatic system for determining the similarity of programs.** That is, your program will be compared with the programs of the other students in class as well as with the programs submitted in previous years. This program will find similarities even if you rename variables, move code, change code structure, etc.

Note that, if you use a Search Engine to find similar programs on the web, we will find these programs too. So don't do it because you will get caught and you will get an F in the course (this is cheating). Don't look for code to use for your project on the web or from other students (current or past). Just do your project alone using the help given in this project description and from your instructor and GTA only.

Platform

You will develop your program on SDSC Expanse. Optionally, you may use your laptop/PC to develop your program first and then, after you make sure that it works there, you transfer it and test it to Expanse. You may also use IntelliJ IDEA or Eclipse to help you develop your program, if you have done so in Project 1. Note that it is required that you test your programs on Expanse before you submit them.

Project Description

For this project, you are asked to implement matrix multiplication in Map-Reduce. This is described in Section 2.3.9 (page 38) in [Map-Reduce and the New Software Stack](#). You should use two Map-Reduce jobs in the same Java file `project2/src/main/java/Multiply.java`. Do not use the method in Section 2.3.10. **You should modify `Multiply.java` only.** In your Java main program, `args[0]` is the first input matrix M, `args[1]` is the second input matrix N, `args[2]` is the directory name to pass the intermediate results from the first Map-Reduce job to the second, and `args[3]` is the output directory. Like in Project 1, all the input and output file formats must be text formats. There are two small sparse matrices 4×3 and 3×3 in the files `M-matrix-small.txt` and `N-matrix-small.txt` for testing in standalone mode. Their matrix multiplication must return the 4×3 matrix in `solution-small.txt`. Then there are 2 moderate-sized matrices 200×100 and 100×300 in the files `M-matrix-large.txt` and `N-matrix-large.txt` for testing in distributed mode. Their matrix multiplication must return the matrix in `solution-large.txt`.

Setting up your Project on your laptop

You can use your laptop to develop your program and then test it and run it on Expanse. Note that testing and running your program on Expanse is required. If you do the project on your laptop, download and untar `project2`:

```
wget http://lambda.uta.edu/cse6331/project2.tgz
tar xzf project2.tgz
```

To compile and run `project2` on your laptop:

```
cd project2
mvn install
rm -rf intermediate output
~/hadoop-3.2.2/bin/hadoop jar target/*.jar Multiply M-matrix-small.txt N-matrix-small.txt intermediate output
```

The file `output/part-r-00000` will contain the results which must be the same as in `solution-small.txt`. After you make sure that the project works correctly for small data on your laptop, copy the files to Expanse and test it on Expanse.

Setting up your Project on Expanse

This step is required. If you'd like, you can develop this project completely on Expanse. If you have already developed `project2` on your laptop, copy `project2.tgz` from your laptop to Expanse. Otherwise, download `project2`:

```
wget http://lambda.uta.edu/cse6331/project2.tgz
tar xzf project2.tgz
```

```
chmod -R g-wrx,o-wrx project2
```

You can compile `Multiply.java` on Expanse using:

```
run multiply.build
```

and you can run it in standalone mode over the two small matrices using:

```
sbatch multiply.local.run
```

The result matrix in the directory output must be similar to `result-matrix-small.txt`. You should modify and run your programs in standalone mode until you get the correct result. After you make sure that your program runs correctly in standalone mode, you run it in distributed mode using:

```
sbatch multiply.distr.run
```

This will multiply the moderate-sized matrices and will write the result in the directory output-distr. Note that running in distributed mode will use up many SUs. So do this once or twice only, after you make sure that your program works correctly in standalone mode.

Pseudo-Code

To help you, I am giving you the pseudo code:

```
class Elem extends Writable {
    short tag; // 0 for M, 1 for N
    int index; // one of the indexes (the other is used as a key)
    double value;
    ...
}

class Pair extends WritableComparable<Pair> {
    int i;
    int j;
    ...
}
```

(Add methods `toString` so you can print `Elem` and `Pair`.) First Map-Reduce job:

```

map(key,line) =           // mapper for matrix M
    split line into 3 values: i, j, and v
    emit(j,new Elem(0,i,v))

map(key,line) =           // mapper for matrix N
    split line into 3 values: i, j, and v
    emit(i,new Elem(1,j,v))

reduce(index,values) =
    A = all v in values with v.tag==0
    B = all v in values with v.tag==1
    for a in A
        for b in B
            emit(new Pair(a.index,b.index),a.value*b.value)

```

Second Map-Reduce job:

```

map(key,value) = // do nothing
    emit(key,value)

reduce(pair,values) = // do the summation
    m = 0
    for v in values
        m = m+v
    emit(pair,m)

```

What to Submit

You need to submit the following files:

```

project2/src/main/java/Multiply.java
project2/multiply.local.out
project2/output-distr/part-r-00000
project2/multiply.distr.out

```

Note: if you cannot get multiply.local.out on Expanse then get it from your laptop/PC by redirecting the output to a file:

```

~/hadoop-3.2.2/bin/hadoop jar target/*.jar Multiply M-matrix-small.txt N-matrix-small.txt intermediate output &>multiply.local.out

```