

# Project 5

New Attempt

**Due** Nov 4 by 11:59pm **Points** 100 **Submitting** a file upload

## Graph Processing on Spark

### Description

The purpose of this project is to develop a graph analysis program using Apache Spark.

This project must be done individually. No copying is permitted. **Note: We will use a system for detecting software plagiarism, called Moss (<http://theory.stanford.edu/~aiken/moss/>), which is an automatic system for determining the similarity of programs.** That is, your program will be compared with the programs of the other students in class as well as with the programs submitted in previous years. This program will find similarities even if you rename variables, move code, change code structure, etc.

Note that, if you use a Search Engine to find similar programs on the web, we will find these programs too. So don't do it because you will get caught and you will get an F in the course (this is cheating). Don't look for code to use for your project on the web or from other students (current or past). Just do your project alone using the help given in this project description and from your instructor and GTA only.

### Update (02/11/2021):

There are many different solutions. So you don't need to return the same results as the solution. The new requirements are that 1) the grand total of the group sums for the small dataset must be 64 and for the large 100000, and 2) that it should not be any -1 group (ie, there should be no unassigned nodes).

### Platform

As in the previous projects, you will develop your program on SDSC Expanse. Optionally, you may use your laptop or IntelliJ Idea or Eclipse to help you develop your program, but you should test your programs on Expanse before you submit them.

## Using your laptop to develop your project

If you'd prefer, you may use your laptop to develop your program and then you would need to test it and run it on Expanse.

To install the project:

```
cd
wget http://lambda.uta.edu/cse6331/project5.tgz
tar xzf project5.tgz
```

To compile and run project5:

```
cd project5
mvn install
~/spark-3.1.2-bin-hadoop3.2/bin/spark-submit --class Partition --master local[2] target/cse6331-project5-0.1.jar small-graph.txt
```

Your result should be the same as the small-result.txt. You may also run it on the large graph:

```
~/spark-3.1.2-bin-hadoop3.2/bin/spark-submit --class Partition --master local[2] target/cse6331-project5-0.1.jar large-graph.txt
```

Your result should be the same as the large-result.txt.

## Setting up your Project on Expanse

Login into Expanse and download and untar project5:

```
wget http://lambda.uta.edu/cse6331/project5.tgz
tar xzf project5.tgz
chmod -R g-wrx,o-wrx project5
```

## Project Description

You are asked to re-implement Project #3 (graph partitioning) using Spark and Scala. **Do not use Map-Reduce.** An empty project5/src/main/scala/Partition.scala is provided, as well as scripts to build and run this code on Expanse. **You should modify Partition.scala only.** Your main program should take the text file that contains the graph (small-graph.txt or large-graph.txt) as an argument.

The graph can be represented as RDD[ ( Long, Long, List[Long] ) ], where the first Long is the graph node ID, the second Long is the assigned cluster ID (-1 if the node has not been assigned yet), and the List[Long] is the adjacent list (the IDs of the neighbors). Here is the code with missing details:

```
var graph = /* read graph from args(0); the graph cluster ID is set to -1 except for the first 5 nodes */

for (i <- 1 to depth)
  graph = graph.flatMap{ /* (1) */ }
                .reduceByKey(_ max _)
                .join( graph.map( /* (2) */ ) )
                .map{ /* (3) */ }

/* finally, print the partition sizes */
```

where the mapper function (1) takes a node ( id, cluster, adjacent) in the graph and returns (id,cluster) along with all (x,cluster) for all x in adjacent. Then the join joins the result with the graph (after it is mapped with function (2)). The join returns an RDD of tuples (id,(new, (old,adjacent))) with the new and the old cluster numbers. In function (3) you keep the old cluster if it's not -1, otherwise you use the new.

Note that in the Expanse script, both the Spark local and distributed modes use 2 RDD partitions. This means that in order to get 10 centroids, we need to get 5 from each partition (the first 5 of each partition).

You can compile Partition.scala using:

```
run partition.build
```

and you can run it in local mode over the small graph using:

```
sbatch partition.local.run
```

Your result should be the same as the small-result.txt. You should modify and run your programs in local mode until you get the correct result. After you make sure that your program runs correctly in local mode, you run it in distributed mode using:

```
sbatch partition.distr.run
```

This will work on the moderate-sized graph and will print the results to the output. It should be the same as large-solution.txt.

## What to Submit

Submit the zipped project5 directory, which must contain the files:

```
project5/src/main/scala/Partition.scala  
project5/partition.local.out  
project5/partition.distr.out
```